

Reading:  
Dominant Graph: An Efficient Indexing Structure to  
Answer Top-K Queries

Yu-Pin Liang

# What is the problem?

## Background

- In a wide range of complex data analytic applications, the query processing often requires computing an expression that involves multiple columns in the relational database. In this study, a broad category of queries which can be expressed as the **scalar product** between a known expression (function) over multiple database attributes and an unknown set of parameters.
- However, in spite of many critical applications of **scalar product query**, surprisingly **no generalized indexing scheme** has been proposed to answer scalar product queries in an online and accurate manner.

# What are the existing techniques and their limitations?

## **Existing techniques and limitations**

Oracle 11.1 release has built-in support for indexing complex SQL functions over multiple attributes. However, their index **does not support** queries that consist of both complex functions as well as unknown parameters.

What makes it possible for the authors to address these limitations?

## **Planar index**

In this work, the problem of fast online computation of scalar product queries in an accurate manner was addressed. To achieve this aim, a novel, lightweight, and generalized indexing scheme, called the **Planar index** was proposed, which can answer the queries very efficiently.

# How is the problem actually solved in this paper?

1. **Offline** technique relies on indexing functions  $\phi(x)$  for data points  $x$  with multiple sets of parallel hyperplanes, and pre-computing some information which is linear in the number of the data points.
2. **Online** query evaluation consists of finding the optimal set of index hyperplanes for a given query, and then using the precomputed information to efficiently answer our queries in an exact manner.
3. The key idea of **Planar** index is to allow very fast pruning of the data points **without actually computing the scalar product** for them.

# Related work

- Half-space range searching
- Linear constraint queries
- Nearest neighbor queries
- Top-k queries with ranking function
- Index for moving objects.

# Problem statement

Given a set of data points in  $\mathbb{R}^d$  and an application specific function  $\phi : \mathbb{R}^d \rightarrow \mathbb{R}^{d'}$ , we define two novel scalar product queries.

**PROBLEM 1 (INEQUALITY QUERY).** *Find all data points  $\mathbf{x} \in \mathbb{R}^d$ , which satisfy a scalar product inequality:  $\langle \mathbf{a}, \phi(\mathbf{x}) \rangle \leq b$ .*

**PROBLEM 2 (TOP- $k$  NEAREST NEIGHBOR QUERY).** *Given some  $k$ , find the top- $k$  data points  $\mathbf{x}$  satisfying  $\langle \mathbf{a}, \phi(\mathbf{x}) \rangle \leq b$ , which also minimize  $\frac{|\langle \mathbf{a}, \phi(\mathbf{x}) \rangle - b|}{|\mathbf{a}|}$ .*

In this study, objective is to propose a generalized indexing scheme—which is easily maintainable and updatable—and which enables faster processing of both the scalar product queries in an accurate manner.

# How is the performance of the proposed technique compared to that of the existing ones?

- Experiments to assess the performance of the **Planar index** for answering scalar product queries were presented.

Table 2: *Dataset characteristics.*

Dataset	# Data Points	# Dimension	Attribute Range
<i>Indp</i>	1,000,000	2 - 14	(1, 100)
<i>Corr</i>	1,000,000	2 - 14	(1, 100)
<i>Anti</i>	1,000,000	2 - 14	(1, 100)
<i>CMoment</i>	68,040	9	(-4.15, 4.59)
<i>CTexture</i>	68,040	16	(-5.25, 50.21)
<i>Consumption</i>	2,075,259	4	(0, 254)

## Competing Method

- Compare the performance of Planar index with a baseline method that performs a naïve sequential scan over the entire dataset.



# How is the performance of the proposed technique compared to that of the existing ones?

**Table 1:** Time complexity of half-space range search algorithms:  $n$  number of data points,  $d$  dimensionality of the query space,  $t$  cardinality of the answer set,  $\epsilon > 0$  any constant,  $c = c(d)$  another constant.

	Query time	Preprocessing storage	Preprocessing time
Agarwal et. al. [1]	$\mathcal{O}(n^{1-\frac{1}{d}+\epsilon} + t)$	$\mathcal{O}(n)$	$\mathcal{O}(n \log n)$
Matousek et. al. [19]	$\mathcal{O}(n^{1-\frac{1}{\lfloor d/2 \rfloor}} (\log n)^c + t)$	$\mathcal{O}(n \log \log n)$	$\mathcal{O}(n \log n)$
Arya et. al. [2]	$\tilde{\Omega}(\frac{n^{1-\frac{1}{d+1}}}{m^{\frac{1}{d+1}}} + t)$	$\tilde{\Omega}(m); \quad n \leq m \leq n^d$	$\mathcal{O}(n^{1+\epsilon} + m(\log n)^\epsilon)$
Planar index [this work]	$\mathcal{O}(d \log n + t) \sim \mathcal{O}(dn)$	$\mathcal{O}(n)$	$\mathcal{O}(n \log n)$

**Table 3:** Top- $k$  nearest-neighbor-finding time using Indp dataset; # dimensions = 6,  $RQ=4$ , # index = 100

# Top- $k$	Checked Points/Total Points (%) [Planar index]	Query Time (ms) [Planar index]	Baseline Time (ms)
50	10.97	33	89
1000	11.29	36	89
10000	12.62	42	89

# How is the performance of the proposed technique compared to that of the existing ones?




Figure 6: Index and query-processing times using real-world datasets (*Consumption*, *CMoment*, and *CTexture*)

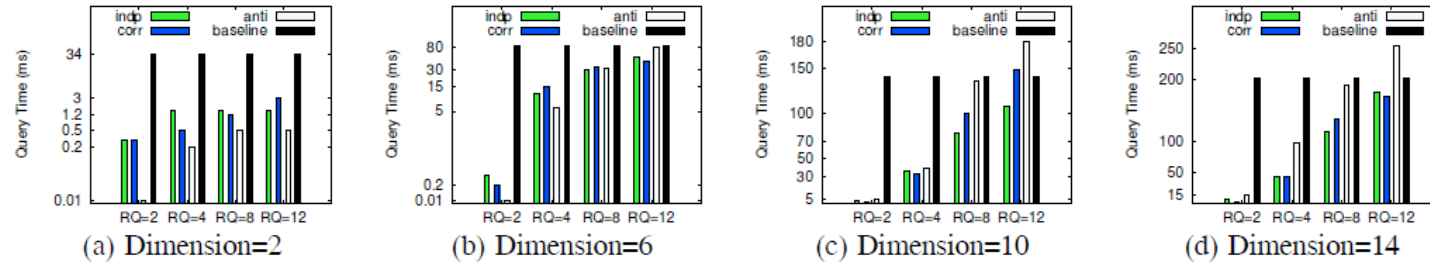


Figure 7: Query-processing time using synthetic datasets (*Indp*, *Corr*, and *Anti*): # dimensions = 2 ~ 14, and randomness of query (RQ) varied from 2 ~ 12, # index = 100. Baseline running times are for any of the three synthetic datasets.




Figure 8: Query-processing time using synthetic datasets (*Indp*, *Corr*, and *Anti*): # dimensions = 2 ~ 14 and # index = 1 ~ 100, randomness of query (RQ) = 4. Baseline running times are for any of the three synthetic datasets.

# How is the performance of the proposed technique compared to that of the existing ones?




Figure 9: Pruning percentage for synthetic datasets: # dimensions = 2 ~ 14, and randomness of query (RQ) = 2 ~ 12, # index = 100.

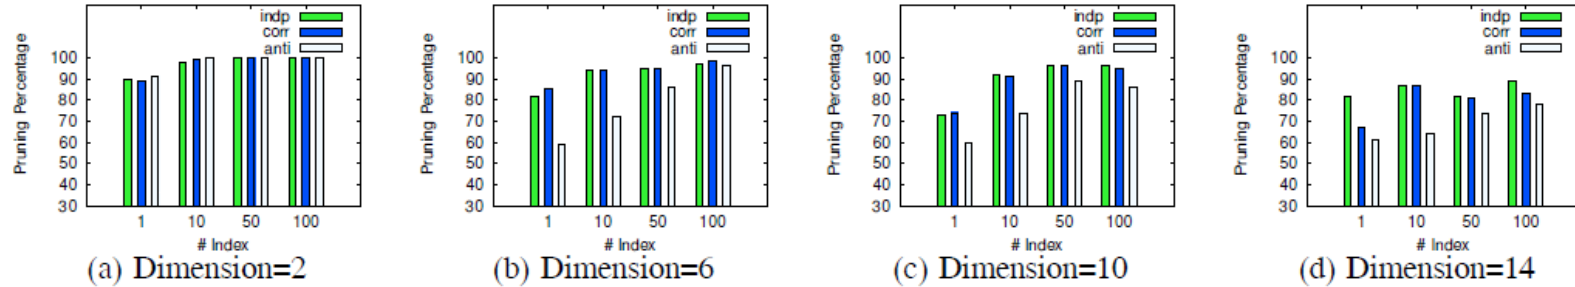


Figure 10: Pruning percentage for synthetic datasets: # dimensions = 2 ~ 14 and # index = 1 ~ 100, randomness of query (RQ) = 4.

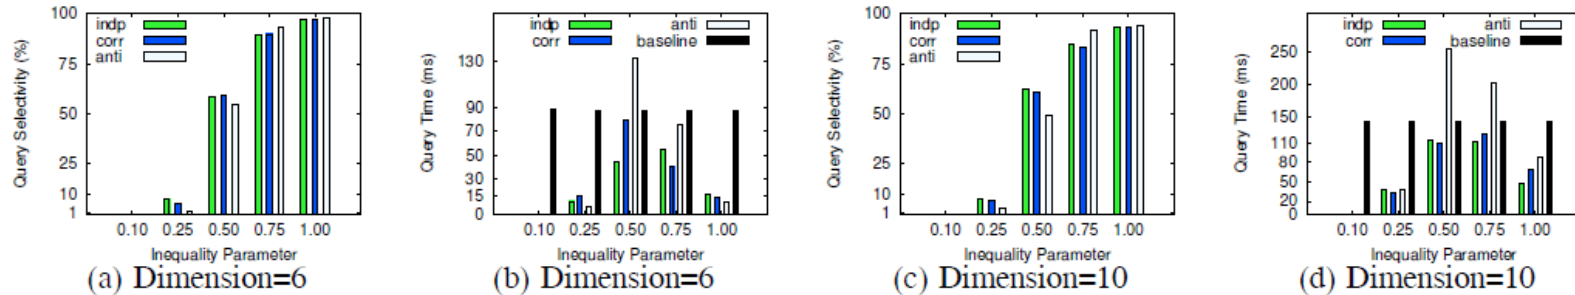


Figure 11: Query selectivity and query-processing time with varying inequality parameter: synthetic datasets, # index = 100, randomness of query (RQ) = 4. Baseline running times are for any of the three synthetic datasets.

# How is the performance of the proposed technique compared to that of the existing ones?

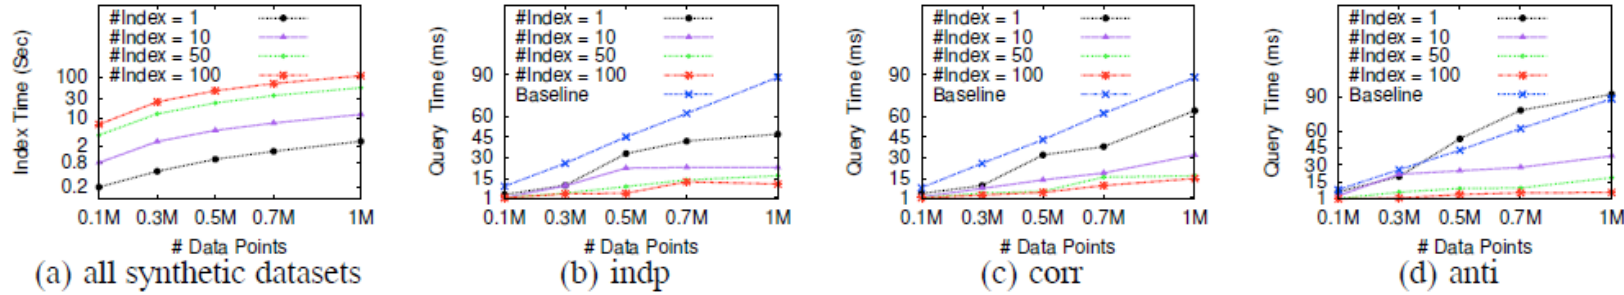


Figure 12: Scalability with varying number of data points using synthetic datasets: # index= 1~100, randomness of query (RQ) = 4, and # dimensions = 6.




Figure 13: Index construction time, memory usage, and dynamic index updates using synthetic datasets




Figure 14: Planar index in finding moving object intersection: uniform and non-uniform workloads

# Conclusion

1. In this paper, scalar product queries was studied — a widely applicable set of analytic queries whose parameters are known only at the time of querying.
2. Defined Planar index, a geometric approach that allows for online processing of scalar product queries in an efficient and accurate manner, as confirmed by an extensive experimental evaluation conducted on various synthetic and real-world datasets.
3. Show the applications of Planar index in the moving-objects-intersection problem and in active learning.

# What do you think about this paper?

- The new indexing scheme 'Planar' for the scalar product queries is innovative and practical for answering online and offline queries.
- Authors also presented a sophisticated approach to compare Planar indexing and normal indexing and make Planar indexing more persuasive.
- However, how to implement the Planar indexing function under the real-world setting is still a concern need to be addressed.

# Reference

- *Arijit Khan, Pouya Yanki, Bojana Dimcheva, Donald Kossmann.* (2014) Towards Indexing Functions: Answering Scalar Product Queries. SIGMOD '14: Proceedings of the 2014 ACM SIGMOD International Conference on Management of Data.